

CUSP-GX-5008.001: Big Data Management & Analysis
SPRING 2016

Homework 7 – Spatial Joins with Apache Spark

In this homework, we would like to generate spatial statistics for yellow taxi trips in NYC. We are interested to know for destinations in each borough of New York, i.e. Manhattan, Brooklyn, Queens, Bronx and Staten Island, the top 3 neighborhoods (e.g. West Village, Williamsburg, Flushing, etc.) that those trips originated from. For example, an answer could be that *Upper East Side*, *Midtown West*, and *Laguardia Airport* are the top 3 origin neighborhood for trips ending up in Manhattan. For the month of May 2011, there were more than 15 million trips, each with pick-up and drop-off location information (i.e. latitude and longitude) and the total size is over 3GB. You are asked to write a Spark application to compute such statistics for the taxi trip records in May 2011. You are also provided with the spatial boundaries for the NYC boroughs and neighborhoods.

DATA SET:

yellow_tripdata_2011-05.csv

Source: http://www.nyc.gov/html/tlc/html/about/trip_record_data.shtml

Description: one-month extract of the TLC Trip Record Data for May 2011. The file is available on HDFS at: `hdfs:///data/share/bdm/hw_final/yellow_tripdata_2011-05.csv`

including with its header.

neighborhoods.geojson

Source: <http://catalog.opendata.city/dataset/pediacities-nyc-neighborhoods>

Description: extracted from the Pediacities NYC Neighborhoods polygons and correlated data, containing only neighborhood geometries, their names and corresponding boroughs. This file is available on HDFS at:

`hdfs:///data/share/bdm/hw_final/neighborhoods.geojson`

boroughs.geojson

Source: <https://data.cityofnewyork.us/City-Government/Borough-Boundaries/tqmj-j8zm>

Description: extracted from NYC Department of City Planning, containing the geometries and name for 5 boroughs of NYC. This file is also available on HDFS at:

`hdfs:///data/share/bdm/hw_final/boroughs.geojson`

Please note that the NYC borough boundaries could be derived from the neighborhoods file as borough information are also included. However, they are both made available for your convenience.

OBJECTIVE: (10 pts)

Please write a Spark application that takes the above file **yellow_tripdata_2011-05.csv** on HDFS as its input and produce the top 3 origin neighborhoods that delivered passengers to each of the five borough based on the number of trips served.

Your submission: you can turn in one or more files including your application's main (Python) file and any dependencies that it may need. However, all of the submitted file(s) must be able to fit into a single **spark-submit** command running on the CUSP's cluster. Please provide this command when submitting your code. For sanity check, please also include the results in your body of the NYU Classes submission.

Evaluation: your code will be tested to run with exactly 22 executors, matching the number of blocks of the input file, where each executor only runs with 1 core. In other words, your code will be run with the following command structure (in a single line):

```
spark-submit --num-executors 22 --py-files ... --packages ... \  
--files "hdfs:///data/share/bdm/hw_final/neighborhoods.geojson,\  
hdfs:///data/share/bdm/hw_final/boroughs.geojson" \  
application.py \  
hdfs:///data/share/bdm/hw_final/yellow_tripdata_2011-05.csv
```

Note: the above command is only an example to demonstrate how to specify the number of executors. It will not run if you just copy and paste into the console. Instead, please check out the **submit6b.sh** script available on NYU Classes as part of the Homework 6 Solutions for further details.

Extra Credit: the top 2 groups with the fastest running times (and correct results) will be given 5 extra points.